

Chapter Nine

Weighted Trees

9.1 Minimum Spanning Tree

9.2 Shortest Path Problems


9.3 Dijkstra's Algorithm

(*9.4 Floyd's Algorithm*)

9.1 Minimum-Cost Spanning Tree

- weighted connected undirected graph
- spanning tree
- cost of spanning tree is sum of edge costs
- find spanning tree that has minimum cost

Example



- Network has 10 edges.
- Spanning tree has only $n - 1 = 7$ edges.
- Need to either select 7 edges or discard 3.

Edge Selection Strategies

- Start with an n -vertex 0 -edge forest.
Consider edges in ascending order of cost.
Select edge if it does not form a cycle together with already selected edges.
 - Kruskal's method.
- Start with a 1 -vertex tree and grow it into an n -vertex tree by repeatedly adding a vertex and an edge. When there is a choice, add a least cost edge.
 - Prim's method.



Edge Selection Strategies

- Start with an **n**-vertex forest. Each component/tree selects a least cost edge to connect to another component/tree. Eliminate duplicate selections and possible cycles. Repeat until only **1** component/tree is left.
 - Sollin's method.

Edge Rejection Strategies



- Start with the connected graph. Repeatedly find a cycle and eliminate the highest cost edge on this cycle. Stop when no cycles remain.
- Consider edges in descending order of cost. Eliminate an edge provided this leaves behind a connected graph.

Kruskal's Method





- Start with a forest that has no edges.
- Consider edges in ascending order of cost.
- Edge (1,2) is considered first and added to the forest.

Kruskal's Method





- Edge (7,8) is considered next and added.
- Edge (3,4) is considered next and added.
- Edge (5,6) is considered next and added.
- Edge (2,3) is considered next and added.
- Edge (1,3) is considered next and rejected because it creates a cycle.

Kruskal's Method





- Edge (2,4) is considered next and rejected because it creates a cycle.
- Edge (3,5) is considered next and added.
- Edge (3,6) is considered next and rejected.
- Edge (5,7) is considered next and added.

Kruskal's Method





- $n - 1$ edges have been selected and no cycle formed.
- So we must have a spanning tree.
- Cost is 46.
- Min-cost spanning tree is unique when all edge costs are different.

(***Prim's Method





- Start with any single vertex tree.
- Get a 2-vertex tree by adding a cheapest edge.
- Get a 3-vertex tree by adding a cheapest edge.
- Grow the tree one edge at a time until the tree has $n - 1$ edges (and hence has all n vertices).

Sollin's Method



- Start with a forest that has no edges.
- Each component selects a least cost edge with which to connect to another component.
- Duplicate selections are eliminated.
- Cycles are possible when the graph has some edges that have the same cost.

Sollin's Method



- Each component that remains selects a least cost edge with which to connect to another component.
- Beware of duplicate selections and cycles.

Minimum-Cost Spanning Tree Methods

- Can prove that all stated edge selection/rejection result in a minimum-cost spanning tree.
- Prim's method is fastest.
 - $O(n^2)$ using an implementation similar to that of Dijkstra's shortest-path algorithm.
 - $O(e + n \log n)$ using a Fibonacci heap.
- Kruskal's uses union-find trees to run in $O(n + e \log e)$ time.

Pseudocode For Kruskal's Method

Start with an empty set T of edges.

while (E is not empty && $|T| \neq n-1$)

{

Let (u,v) be a least-cost edge in E .

$E = E - \{(u,v)\}$. // delete edge from E

if ((u,v) does not create a cycle in T)

 Add edge (u,v) to T .

}

if ($|T| == n-1$) T is a min-cost spanning tree.

else Network has no spanning tree.

Data Structures For Kruskal's Method

Edge set E .

Operations are:

- Is E empty?
- Select and remove a least-cost edge.

Use a min heap of edges.

- Initialize. $O(e)$ time.
- Remove and return least-cost edge. $O(\log e)$ time.

Data Structures For Kruskal's Method

Set of selected edges T .

Operations are:

- Does T have $n - 1$ edges?
- Does the addition of an edge (u, v) to T result in a cycle?
- Add an edge to T .


Data Structures For Kruskal's Method

Use an array for the edges of **T**.

- Does **T** have **$n - 1$** edges?
 - Check number of edges in array. **$O(1)$** time.
- Does the addition of an edge **(u, v)** to **T** result in a cycle?
 - Not easy.
- Add an edge to **T**.
 - Add at right end of edges in array. **$O(1)$** time.


Data Structures For Kruskal's Method

Does the addition of an edge (u, v) to T result in a cycle?



- Each component of T is a tree.
- When u and v are in the same component, the addition of the edge (u, v) creates a cycle.
- When u and v are in the different components, the addition of the edge (u, v) does not create a cycle.


Data Structures For Kruskal's Method



- Each component of T is defined by the vertices in the component.
- Represent each component as a set of vertices.
 - $\{1, 2, 3, 4\}, \{5, 6\}, \{7, 8\}$
- Two vertices are in the same component iff they are in the same set of vertices.

Data Structures For Kruskal's Method

- Initially, **T** is empty.



- Initial sets are:
 - $\{1\} \{2\} \{3\} \{4\} \{5\} \{6\} \{7\} \{8\}$
- Does the addition of an edge (u, v) to **T** result in a cycle? If not, add edge to **T**.
 $s1 = \text{Find}(u); s2 = \text{Find}(v);$
 $\text{if } (s1 \neq s2) \text{ Union}(s1, s2);$


Data Structures For Kruskal's Method

- Use fast solution for disjoint sets.
- Initialize.
 - $O(n)$ time.
- At most $2e$ finds and $n-1$ unions.
 - Very close to $O(n + e)$.
- Min heap operations to get edges in increasing order of cost take $O(e \log e)$.
- Overall complexity of Kruskal's method is $O(n + e \log e)$.

9.2 Shortest Path Problems


- Directed weighted graph.
- Path length is sum of weights of edges on path.
- The vertex at which the path begins is the **source** vertex.
- The vertex at which the path ends is the **destination** vertex.

Example



A path from 1 to 7.
Path length is 14.

Example



Another path from 1 to 7.
Path length is 11.

Shortest Path Problems


- Single source single destination.
- Single source all destinations.
- All pairs (every vertex is a source and destination).

Single Source Single Destination

Possible algorithm: (*****)

- Leave source vertex using cheapest/shortest edge.
- Leave new vertex using cheapest edge subject to the constraint that a new vertex is reached.
- Continue until destination is reached.

Constructed 1 To 7 Path



Path length is 12.

Not shortest path. Algorithm doesn't work!


Single Source All Destinations

Need to generate up to n (n is number of vertices) paths (including path from source to itself).

Dijkstra's method:

- Construct these up to n paths in order of increasing length.
- Assume edge costs (lengths) are ≥ 0 .
- So, no path has length < 0 .
- First shortest path is from the source vertex to itself. The length of this path is 0 .

Single Source All Destinations




Path	Length		
①	0	① → ②	6
① → ③	2	① → ③ → ⑤ → ④	9
① → ③ → ⑤	5	① → ③ → ⑥	10
		① → ③ → ⑥ → ⑦	11

Single Source All Destinations


Path Length




0




2




5




6



9



10



11


- Each path (other than first) is a one edge extension of a previous path.

- Next shortest path is the shortest one edge extension of an already generated shortest path.



Single Source All Destinations

- Let $d[i]$ be the length of a shortest one edge extension of an already generated shortest path, the one edge extension ends at vertex i .
- The next shortest path is to an as yet unreached vertex for which the $d[]$ value is least.
- Let $p[i]$ be the vertex just before vertex i on the shortest one edge extension to i .

Single Source All Destinations





Single Source All Destinations





	[1]	[2]	[3]	[4]	[5]	[6]	[7]
d	0	6	2	16	5	10	14
p	-	1	1	1	3	3	1

Single Source All Destinations





	[1]	[2]	[3]	[4]	[5]	[6]	[7]
d	0	6	2	9	5	10	14
p	-	1	1	5	3	3	1

Single Source All Destinations





	[1]	[2]	[3]	[4]	[5]	[6]	[7]
d	0	6	2	9	5	10	14
p	-	1	1	5	3	3	1

Single Source All Destinations



	[1]	[2]	[3]	[4]	[5]	[6]	[7]
d	0	6	2	9	5	10	12
p	-	1	1	5	3	3	4

Single Source All Destinations



	[1]	[2]	[3]	[4]	[5]	[6]	[7]
d	0	6	2	9	5	10	11
p	-	1	1	5	3	3	6

Single Source All Destinations

Path	Length	[1]	[2]	[3]	[4]	[5]	[6]	[7]
1	0	0	-	-	-	-	-	-
1 → 3	2	6	1	-	-	-	-	-
1 → 3 → 5	5	2	1	5	3	-	-	-
1 → 2	6	9	-	-	-	-	-	-
1 → 3 → 5 → 4	9	5	3	3	6	-	-	-
1 → 3 → 6	10	-	-	-	-	-	-	-
1 → 3 → 6 → 7	11	-	-	-	-	-	-	-

Single Source Single Destination

Terminate single source all destinations algorithm as soon as shortest path to desired vertex has been generated.

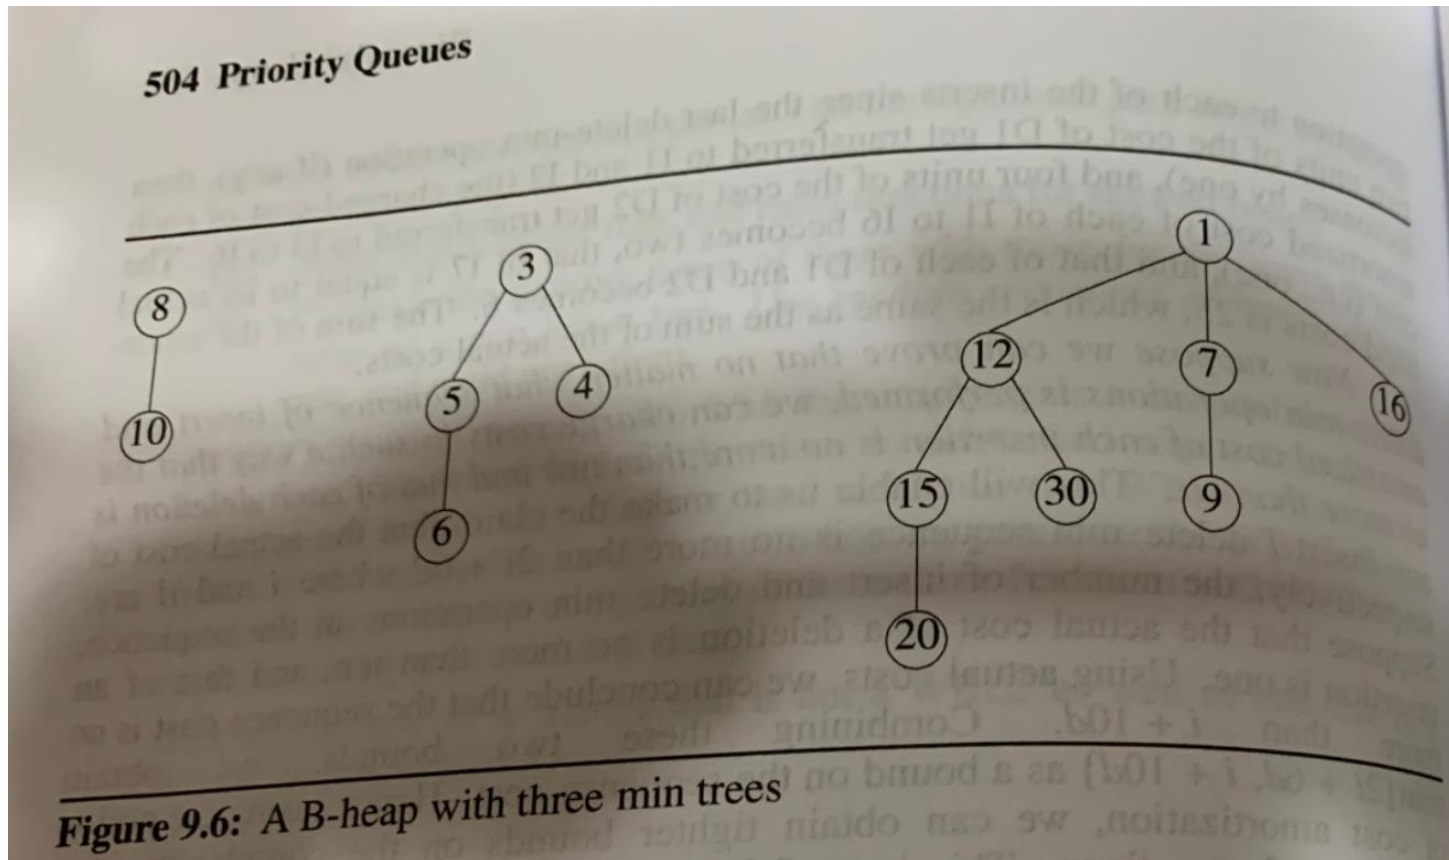
Fibonacci Heaps

- **Def:** There are two varieties of Fibonacci heaps: min and max. A min Fibonacci heap is a collection of min trees; a max Fibonacci heap is a collection of max trees.
- Here considered is min Fibonacci heap, and named as F_heaps.
- An F-heap is a data structure that support the seven operations: getMin, Insert, DeleteMin, Meld, Delete, and DecreaseKey, where the last two are defined as follows:
 - **Delete:** Delete the element in a specified node. We refer too this delete operation as arbitrary delete.
 - **DecreaseKey:** Decrease the key/priority of a specified node by a given positive amount.

Fibonacci Heaps(c.)

- When an F-heap, is used, the Delete operation takes $O(\log n)$ amortized time and the DecreaseKey takes $O(1)$ amortized time.
- During implementation, there are at least five pointers inside each node: parent, child, childCut, leftlink and rightlink, one datum representing the number of its children, and one datum associated with a key. The children nodes of a node can be connected by a double linked list according to leftlink and rightlink correspondingly.

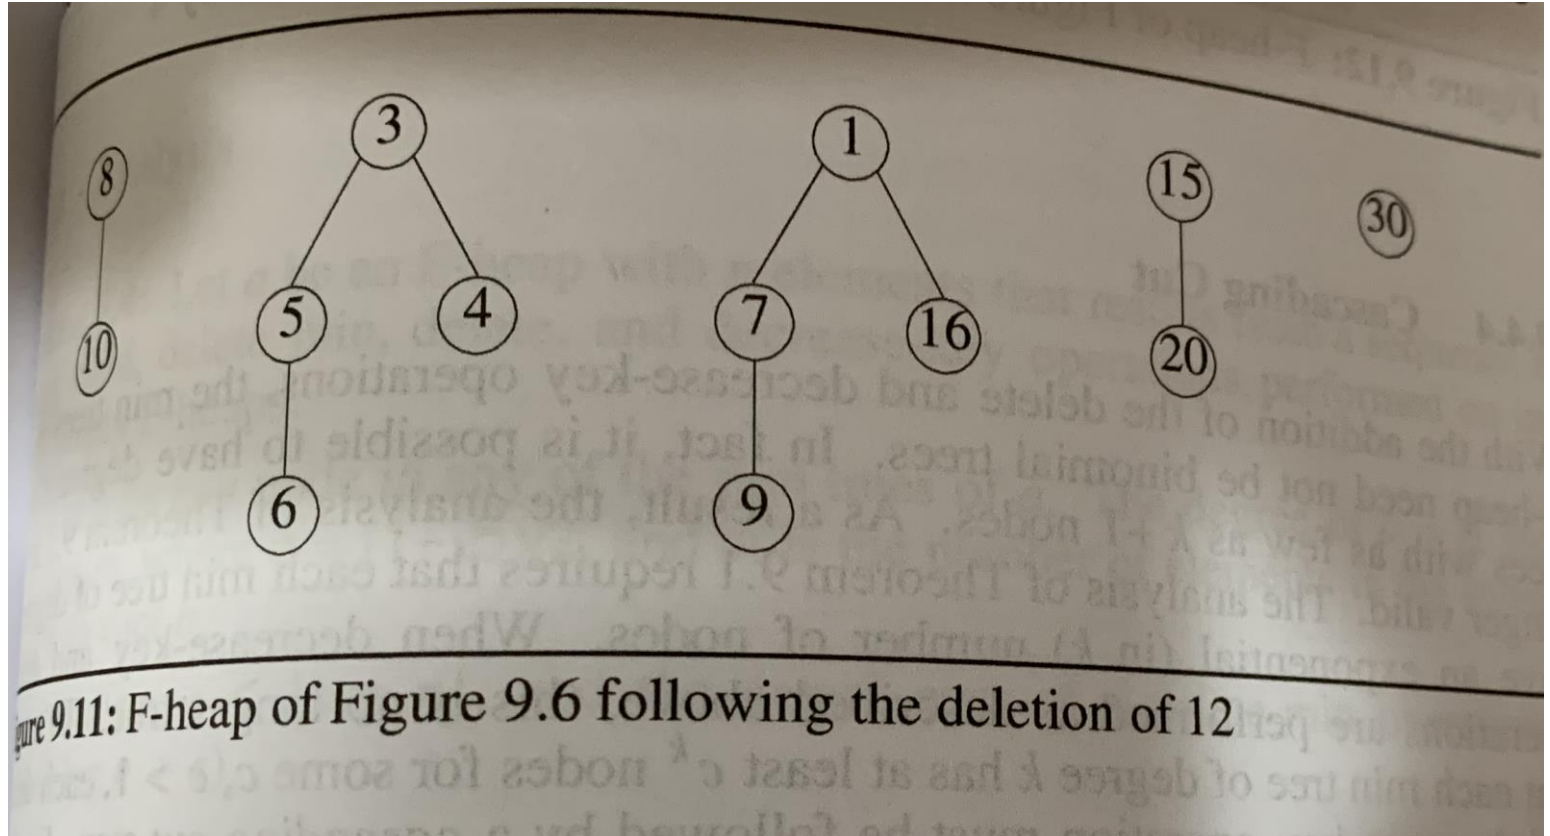
Fibonacci Heaps(c.)



Fibonacci Heaps(c.)

- Deletion from an F-Heap
- To delete an arbitrary node b from a F-heap, we do the following:
 1. If $\text{min}=b$, then do a delete-min; otherwise do Steps 2, 3, and 4 below.
 2. Delete b from its doubly linked list.
 3. Combine the doubly linked list of b 's children with the doubly linked list pointed at by min into a single double linked list. Trees of equal degree are not jointed as a delete-min operation.
 4. Dispose of node b .


Fibonacci Heaps(c.)



Fibonacci Heaps(c.)

- Decrease Key
- To decrease the key in node b we do the following:
 1. Reduce the key in b .
 2. If b is not a min tree root and its key is smaller than that in its parent, then delete b from its doubly linked list and insert it into the doubly linked list of min tree nodes
 3. Change min to point to b if the key in b is smaller than that in min.

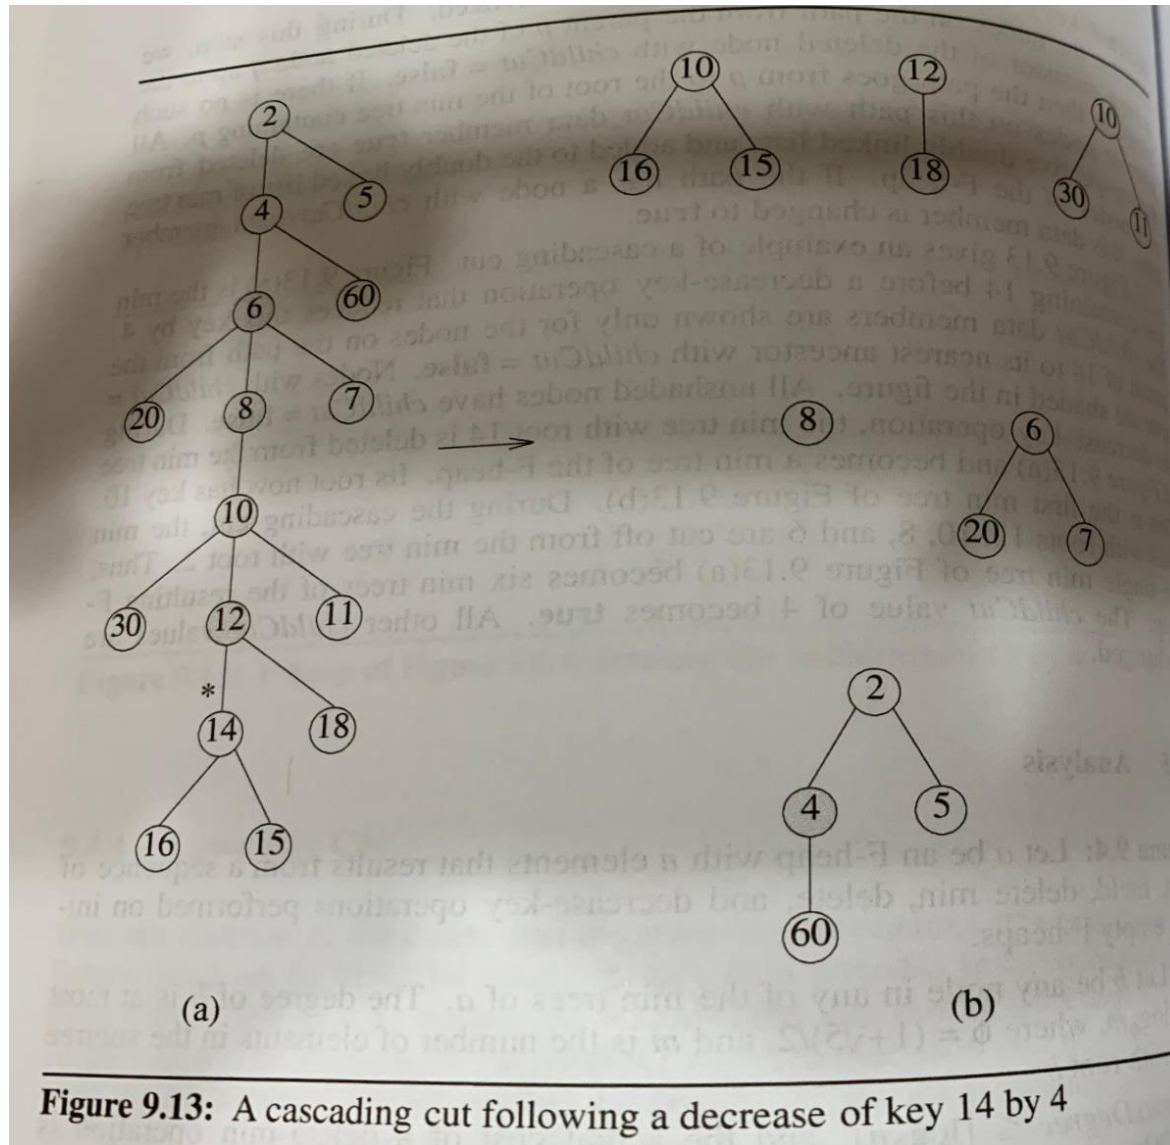
Fibonacci Heaps(c.)



Fibonacci Heaps(c.)

- Cascading Cut
- To decrease the key in node b we do the following:
 1. Reduce the key in b .
 2. If b is not a min tree root and its key is smaller than that in its parent, then delete b from its doubly linked list and insert it into the doubly linked list of min tree nodes
 3. Change min to point to b if the key in b is smaller than that in min .

Fibonacci Heaps(c.)



Fibonacci Heaps(c.)

Analysis :

- Lemma 1: Let a be an F-heap with n elements that results from a sequence of insert, meld, delete min, delete, and decrease-key operations performed on initial empty F-heaps.
 1. Let b be any node in any of the min trees of a . The degree of b is at most $\log_{\phi} m$, where $\phi = (1 + \sqrt{5})/2$, and m is the number of elements in the subtree with root b .
 2. $\text{maxDegree} \leq \lceil \log_{\phi} n \rceil$, and the actual cost of a delete-min operation is $O(\log n + s)$
- Theorem 2: If a sequence of n insert, meld, delete min, delete, and decrease-key operations is performed on an initially empty F-heap, then we can amortize costs such that the amortized time complexity of each insert, meld, and decrease-key operation is $O(1)$ and that of each delete min and delete operation is $O(\log n)$. The total time complexity of the entire sequence is the sum of the amortized complexities of the individual operations in the sequence.

9.3 Dijkstra's Algorithm

Data Structures for Dijkstra's Algorithm

- The described single source all destinations algorithm is known as Dijkstra's algorithm.
- Implement $d[]$ and $p[]$ as 1D arrays.
- Keep a linear list L of reachable vertices to which shortest path is yet to be generated.
- Select and remove vertex v in L that has smallest $d[]$ value.
- Update $d[]$ and $p[]$ values of vertices adjacent to v .

Complexity



- $O(n)$ to select next destination vertex.
- $O(\text{out-degree})$ to update $d[]$ and $p[]$ values when adjacency lists are used.
- $O(n)$ to update $d[]$ and $p[]$ values when adjacency matrix is used.
- Selection and update done once for each vertex to which a shortest path is found.
- Total time is $O(n^2 + e) = O(n^2)$.

Complexity



- When a min heap of $d[]$ values is used in place of the linear list L of reachable vertices, total time is $O((n+e) \log n)$, because $O(n)$ remove min operations and $O(e)$ change key ($d[]$ value) operations are done.
- When e is $O(n^2)$, using a min heap is worse than using a linear list.
- When a Fibonacci heap is used, the total time is $O(n \log n + e)$.

Single-Source All-Destinations Shortest Paths With General Weights

- Directed weighted graph.
- Edges may have negative cost.
- No cycle whose cost is < 0 .
- Find a shortest path from a given source vertex s to each of the n vertices of the digraph.

Single-Source All-Destinations Shortest Paths With General Weights

- Dijkstra's $O(n^2)$ single-source greedy algorithm doesn't work when there are negative-cost edges.

Bellman-Ford Algorithm

- Single-source all-destinations shortest paths in digraphs with negative-cost edges.
- Uses dynamic programming.
- Runs in $O(n^3)$ time when adjacency matrices are used.
- Runs in $O(ne)$ time when adjacency lists are used.

Strategy



- To construct a shortest path from the source to vertex **v**, decide on the max number of edges on the path and on the vertex that comes just before **v**.
- Since the digraph has no cycle whose length is < 0 , we may limit ourselves to the discovery of cycle-free (acyclic) shortest paths.
- A path that has no cycle has at most **$n-1$** edges.

Cost Function d



- Let $d(v, k)$ ($\text{dist}^k[v]$) be the length of a shortest path from the source vertex to vertex v under the constraint that the path has at most k edges.
- $d(v, n-1)$ is the length of a shortest unconstrained path from the source vertex to vertex v .
- We want to determine $d(v, n-1)$ for every vertex v .

Value Of $d(*,0)$

- $d(v,0)$ is the length of a shortest path from the source vertex to vertex v under the constraint that the path has at most 0 edges.



- $d(s,0) = 0$.
- $d(v,0) = \text{infinity}$ for $v \neq s$.

Recurrence For $d(*,k)$, $k > 0$

- $d(v,k)$ is the length of a shortest path from the source vertex to vertex v under the constraint that the path has at most k edges.
- If this constrained shortest path goes through no more than $k-1$ edges, then $d(v,k) = d(v,k-1)$.

Recurrence For $d(*,k)$, $k > 0$

- If this constrained shortest path goes through k edges, then let w be the vertex just before v on this shortest path (note that w may be s).



- We see that the path from the source to w must be a shortest path from the source vertex to vertex w under the constraint that this path has at most $k-1$ edges.
- $d(v,k) = d(w,k-1) + \text{length of edge } (w,v)$.

Recurrence For $d(*,k)$, $k > 0$

- $d(v,k) = d(w,k-1) + \text{length of edge } (w,v)$.



- We do not know what w is.
- We can assert
 - $d(v,k) = \min\{d(w,k-1) + \text{length of edge } (w,v)\}$, where the \min is taken over all w such that (w,v) is an edge of the digraph.
- Combining the two cases considered yields:
 - $d(v,k) = \min\{d(v,k-1), \min\{d(w,k-1) + \text{length of edge } (w,v)\}\}$

Pseudocode To Compute $d(*,*)$

// initialize $d(*,0)$

$d(s,0) = 0;$

$d(v,0) = \text{infinity}, v \neq s;$

// compute $d(*,k), 0 < k < n$

for (int $k = 1; k < n; k++$)

{

$d(v,k) = d(v,k-1), 1 \leq v \leq n;$

for (each edge (u,v))


$d(v,k) = \min\{d(v,k), d(u,k-1) + \text{cost}(u,v)\}$

}

$$p(*,*)$$

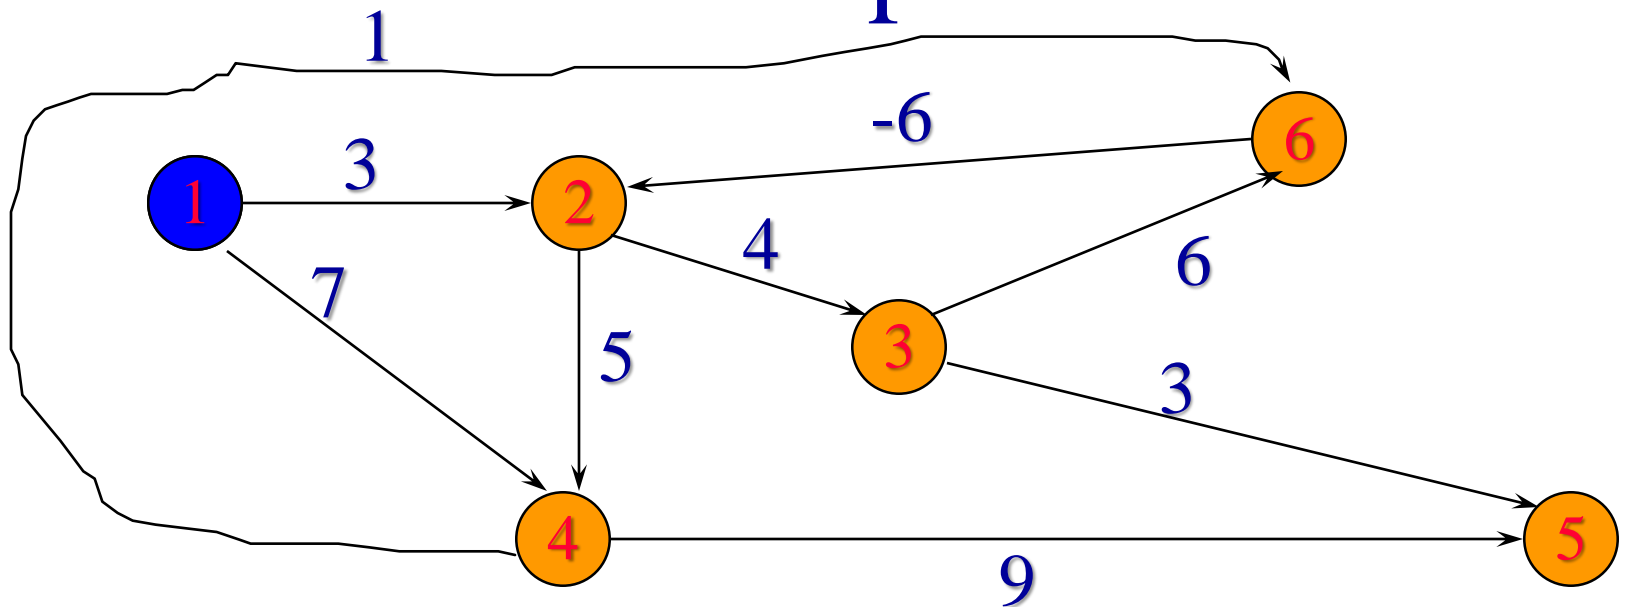
- Let $p(v,k)$ be the vertex just before vertex v on the shortest path for $d(v,k)$.
- $p(v,0)$ is undefined.
- Used to construct shortest paths.

Example



Source vertex is 1.

Example




	1	2	3	4	5	6
0	0	-	-	-	-	-
1	0	3	-	7	-	-
2	0	3	7	7	16	8
3	0	2	7	7	10	8
4	0	2	6	7	10	8

$d(v,k)$

$v \longrightarrow$						
	-	-	-	-	-	-
	-	1	-	1	-	-
	-	1	2	1	4	4
	-	6	2	1	3	4
	-	6	2	1	3	4

$p(v,k)$

Example



	1	2	3	4	5	6
4	0	2	6	7	10	8
5	0	2	6	7	9	8


$d(v,k)$

$v \longrightarrow$					
-	6	2	1	3	4
-	6	2	1	3	4

$k \downarrow$

$p(v,k)$

Shortest Path From 1 To 5



	1	2	3	4	5	6
5	0	2	6	7	9	8

$d(v,5)$

	1	2	3	4	5	6
	-	6	2	1	3	4

$p(v,5)$

Observations

- $d(v,k) = \min\{d(v,k-1), \min\{d(w,k-1) + \text{length of edge } (w,v)\}\}$
- $d(s,k) = 0$ for all k .
- If $d(v,k) = d(v,k-1)$ for all v , then $d(v,j) = d(v,k-1)$, for all $j \geq k-1$ and all v .
- If we stop computing as soon as we have a $d(*,k)$ that is identical to $d(*,k-1)$ the run time becomes
 - $O(n^3)$ when adjacency matrix is used.
 - $O(ne)$ when adjacency lists are used.

Observations

- The computation may be done in-place.

$$d(v) = \min\{d(v), \min\{d(w) + \text{length of edge } (w,v)\}\}$$


instead of

$$d(v,k) = \min\{d(v,k-1), \\ \min\{d(w,k-1) + \text{length of edge } (w,v)\}\}$$

- Following iteration k , $d(v,k+1) \leq d(v) \leq d(v,k)$
- On termination $d(v) = d(v,n-1)$.
- Space requirement becomes $O(n)$ for $d(*)$ and $p(*)$.


• All-Pairs Shortest Paths •

- Given an n -vertex directed weighted graph, find a shortest path from vertex i to vertex j for each of the n^2 vertex pairs (i,j) .




Dijkstra's Single Source Algorithm

- Use Dijkstra's algorithm **n** times, once with each of the **n** vertices as the source vertex.



Performance




- Time complexity is $O(n^3)$ time.
- Works only when no edge has a cost < 0 .

(** 9.4 Floyd's Algorithm **)


- Time complexity is $\Theta(n^3)$ time.
- Works so long as there is no cycle whose length is < 0 .
- When there is a cycle whose length is < 0 , some shortest paths aren't finite.
 - If vertex **1** is on a cycle whose length is -2 , each time you go around this cycle once you get a **1** to **1** path that is **2** units shorter than the previous one.
- Simpler to code, smaller overheads.

Decision Sequence




- First decide the highest intermediate vertex (i.e., largest vertex number) on the shortest path from *i* to *j*.
- If the shortest path is *i*, 2, 6, 3, 8, 5, 7, *j* the first decision is that vertex 8 is an intermediate vertex on the shortest path and no intermediate vertex is larger than 8.
- Then decide the highest intermediate vertex on the path from *i* to 8, and so on.

A Triple



- (i, j, k) denotes the problem of finding the shortest path from vertex i to vertex j that has no intermediate vertex larger than k .
- (i, j, n) denotes the problem of finding the shortest path from vertex i to vertex j (with no restrictions on intermediate vertices).


Cost Function



- Let $c(i,j,k)$ be the length of a shortest path from vertex i to vertex j that has no intermediate vertex larger than k .


$c(i,j,n)$

- $c(i,j,n)$ is the length of a shortest path from vertex i to vertex j that has no intermediate vertex larger than n .
- No vertex is larger than n .
- Therefore, $c(i,j,n)$ is the length of a shortest path from vertex i to vertex j .




$$c(i,j,0)$$

- $c(i,j,0)$ is the length of a shortest path from vertex i to vertex j that has no intermediate vertex larger than 0 .
 - Every vertex is larger than 0 .
 - Therefore, $c(i,j,0)$ is the length of a single-edge path from vertex i to vertex j .




Recurrence For $c(i,j,k)$, $k > 0$

- The shortest path from vertex i to vertex j that has no intermediate vertex larger than k may or may not go through vertex k .
- If this shortest path does not go through vertex k , the largest permissible intermediate vertex is $k-1$. So the path length is $c(i,j,k-1)$.




Recurrence For $c(i,j,k)$, $k > 0$

- Shortest path goes through vertex k .




- We may assume that vertex k is not repeated because no cycle has negative length.
- Largest permissible intermediate vertex on i to k and k to j paths is $k-1$.

Recurrence For $c(i,j,k)$, $k > 0$




- i to k path must be a shortest i to k path that goes through no vertex larger than $k-1$.
- If not, replace current i to k path with a shorter i to k path to get an even shorter i to j path.

Recurrence For $c(i,j,k)$, $k > 0$



- Similarly, k to j path must be a shortest k to j path that goes through no vertex larger than $k-1$.
- Therefore, length of i to k path is $c(i,k,k-1)$, and length of k to j path is $c(k,j,k-1)$.
- So, $c(i,j,k) = c(i,k,k-1) + c(k,j,k-1)$.

Recurrence For $c(i,j,k)$, $k > 0$



- Combining the two equations for $c(i,j,k)$, we get $c(i,j,k) = \min\{c(i,j,k-1), c(i,k,k-1) + c(k,j,k-1)\}$.
- We may compute the $c(i,j,k)$ s in the order $k = 1, 2, 3, \dots, n$.

Floyd's Shortest Paths Algorithm


```
for (int k = 1; k <= n; k++)  
    for (int i = 1; i <= n; i++)  
        for (int j = 1; j <= n; j++)  
            c(i,j,k) = min{ c(i,j,k-1),  
                           c(i,k,k-1) + c(k,j,k-1) };
```

- Time complexity is $O(n^3)$.
- More precisely $\Theta(n^3)$.
- $\Theta(n^3)$ space is needed for $c(*,*,*)$.



Space Reduction

- $c(i,j,k) = \min\{c(i,j,k-1), c(i,k,k-1) + c(k,j,k-1)\}$
- When neither i nor j equals k , $c(i,j,k-1)$ is used only in the computation of $c(i,j,k)$.
column k



- So $c(i,j,k)$ can overwrite $c(i,j,k-1)$.

Space Reduction

- $c(i,j,k) = \min\{c(i,j,k-1), c(i,k,k-1) + c(k,j,k-1)\}$
- When i equals k , $c(i,j,k-1)$ equals $c(i,j,k)$.
 - $c(k,j,k) = \min\{c(k,j,k-1), c(k,k,k-1) + c(k,j,k-1)\}$
 $= \min\{c(k,j,k-1), 0 + c(k,j,k-1)\}$
 $= c(k,j,k-1)$
- So, when i equals k , $c(i,j,k)$ can overwrite $c(i,j,k-1)$.
- Similarly when j equals k , $c(i,j,k)$ can overwrite $c(i,j,k-1)$.
- So, in all cases $c(i,j,k)$ can overwrite $c(i,j,k-1)$.

Floyd's Shortest Paths Algorithm

```
for (int k = 1; k <= n; k++)  
    for (int i = 1; i <= n; i++)  
        for (int j = 1; j <= n; j++)  
             $c(i,j) = \min\{c(i,j), c(i,k) + c(k,j)\};$ 
```

- Initially, $c(i,j) = c(i,j,0)$.
- Upon termination, $c(i,j) = c(i,j,n)$.
- Time complexity is $\Theta(n^3)$.
- $\Theta(n^2)$ space is needed for $c(*,*)$.




Building The Shortest Paths

- Let $\text{kay}(i,j)$ be the largest vertex on the shortest path from i to j .
- Initially, $\text{kay}(i,j) = 0$ (shortest path has no intermediate vertex).

```
for (int k = 1; k <= n; k++)  
    for (int i = 1; i <= n; i++)  
        for (int j = 1; j <= n; j++)  
            if (c(i,j) > c(i,k) + c(k,j))  
                {kay(i,j) = k; c(i,j) = c(i,k) + c(k,j);}
```

Example



-	7	5	1	-	-	-	-
-	-	-	-	4	-	-	-
-	7	-	-	9	9	-	-
-	5	-	-	-	-	16	-
-	-	-	4	-	-	-	1
-	-	-	-	-	-	1	-
2	-	-	-	-	-	-	4
-	-	-	-	-	2	4	-

Initial Cost Matrix

$$c(*,*) = c(*,*,0)$$

Final Cost Matrix $c(*,*) = c(*,*,n)$

0	6	5	1	10	13	14	11
10	0	15	8	4	7	8	5
12	7	0	13	9	9	10	10
15	5	20	0	9	12	13	10
6	9	11	4	0	3	4	1
3	9	8	4	13	0	1	5
2	8	7	3	12	6	0	4
5	11	10	6	15	2	3	0

kay Matrix

0 4 0 0 4 8 8 5

8 0 8 5 0 8 8 5

7 0 0 5 0 0 6 5

8 0 8 0 2 8 8 5


8 4 8 0 0 8 8 0

7 7 7 7 7 0 0 7

0 4 1 1 4 8 0 0

7 7 7 7 7 0 6 0

Shortest Path



Shortest path from 1 to 7.

Path length is 14.

Build A Shortest Path

0 4 0 0 4 8 8 5

8 0 8 5 0 8 8 5

7 0 0 5 0 0 6 5

8 0 8 0 2 8 8 5

8 4 8 0 0 8 8 0

7 7 7 7 7 0 0 7

0 4 1 1 4 8 0 0

7 7 7 7 7 0 6 0

- The path is 1 4 2 5 8 6 7.

- $\text{kay}(1,7) = 8$

1 \longrightarrow 8 \longrightarrow 7

- $\text{kay}(1,8) = 5$

1 \longrightarrow 5 \longrightarrow 8 \longrightarrow 7

- $\text{kay}(1,5) = 4$

1 \longrightarrow 4 \longrightarrow 5 \longrightarrow 8 \longrightarrow 7

Build A Shortest Path

0 4 0 0 4 8 8 5

8 0 8 5 0 8 8 5

7 0 0 5 0 0 6 5

8 0 8 0 2 8 8 5

8 4 8 0 0 8 8 0

7 7 7 7 7 0 0 7

0 4 1 1 4 8 0 0

7 7 7 7 7 0 6 0

- The path is 1 4 2 5 8 6 7.

1 → 4 → 5 → 8 → 7

- $\text{kay}(1,4) = 0$

1 4 → 5 → 8 → 7

- $\text{kay}(4,5) = 2$

1 4 → 2 → 5 → 8 → 7

- $\text{kay}(4,2) = 0$

1 4 2 → 5 → 8 → 7

Build A Shortest Path

0 4 0 0 4 8 8 5

8 0 8 5 0 8 8 5

7 0 0 5 0 0 6 5

8 0 8 0 2 8 8 5

8 4 8 0 0 8 8 0

7 7 7 7 7 0 0 7

0 4 1 1 4 8 0 0

7 7 7 7 7 0 6 0

- The path is 1 4 2 5 8 6 7.

1 4 2 \rightarrow 5 \rightarrow 8 \rightarrow 7

- $\text{kay}(2,5) = 0$

1 4 2 5 \rightarrow 8 \rightarrow 7

- $\text{kay}(5,8) = 0$

1 4 2 5 8 \rightarrow 7

- $\text{kay}(8,7) = 6$

1 4 2 5 8 \rightarrow 6 \rightarrow 7

Build A Shortest Path

0 4 0 0 4 8 8 5

8 0 8 5 0 8 8 5

7 0 0 5 0 0 6 5

8 0 8 0 2 8 8 5

8 4 8 0 0 8 8 0

7 7 7 7 7 0 0 7

0 4 1 1 4 8 0 0

7 7 7 7 7 0 6 0

- The path is 1 4 2 5 8 6 7.

1 4 2 5 8 \rightarrow 6 \rightarrow 7

- $\text{kay}(8,6) = 0$

1 4 2 5 8 6 \rightarrow 7

- $\text{kay}(6,7) = 0$

1 4 2 5 8 6 7

Output A Shortest Path

```
void outputPath(int i, int j)
{ // does not output first vertex (i) on path
  if (i == j) return;
  if (kay[i][j] == 0) // no intermediate vertices on path
    cout << j << " ";
  else { // kay[i][j] is an intermediate vertex on the path
    outputPath(i, kay[i][j]);
    outputPath(kay[i][j], j);
  }
}
```

Time Complexity Of outputPath

$O(\text{number of vertices on shortest path})$